

AI CRITIQUE DOCUMENT

Task A: AI Test Design with Critique
Digital Wallet Transfer Limits & Account Tiers

PREPARED BY	AI TOOL	STATUS
Jeremiah Igbiniedion	ChatGPT	Final

76 AI-GENERATED CASES | 38 REMOVED OR MERGED | 38 FOCUSED CASES

1. Purpose and Method

ChatGPT was used to generate an initial test suite from the supplied Digital Wallet Transfer Limits & Account Tiers specification. The goal was to assess how effectively an LLM could identify functional coverage, boundaries, negative scenarios and edge cases. The output was not accepted automatically. Every case was manually compared with the seven stated requirements and classified as retained, corrected, redundant, merged or unsupported.

REVIEW PRINCIPLE

A useful test case had to be traceable to the supplied specification. A technically interesting scenario was not treated as valid merely because it could occur in a real payment system.

2. Prompts Used

Prompt 1: Initial comprehensive generation

Generate comprehensive test cases for the following feature specification.

Feature: Digital Wallet Transfer Limits & Account Tiers

Requirements:

1. Tier 1: 20,000 per transaction and 50,000 per day. Tier 2: 200,000 per transaction and 500,000 per day. Tier 3: no per-transaction limit and 5,000,000 per day.
2. Daily limit resets at midnight.
3. Pending transfers count toward the daily limit. Failed or reversed amounts are returned to the available daily limit.
4. Transfers to the user's own linked bank accounts are exempt from all limits.
5. Future transfers execute at 8:00 AM on the scheduled day.
6. KYC tier upgrades take effect immediately.
7. Transfers exceeding the remaining limit are rejected with a clear message showing the remaining limit.

Generate detailed cases with ID, priority, requirement, preconditions, test data, steps and expected results.

Prompt 2: Human-led audit and correction

Audit all 76 generated cases against only the seven stated requirements. Identify duplicates, unsupported assumptions, hallucinated product capabilities and incorrect expected results. Keep the defensible cases, explain every removal, and correct the cross-day failure, own-account exemption, scheduled-transfer and KYC-upgrade cases.

3. Initial AI Output

The model produced 76 cases across tier limits, daily resets, transfer status changes, own-account exemptions, scheduling, KYC upgrades, error messages, APIs, concurrency, security and data integrity. The breadth was useful, but the suite also contained repeated checks, implementation assumptions and behaviours absent from the specification.

4. Cases Kept

The following 32 cases were retained without major correction because they directly test stated behaviour. Six additional cases were retained after correction or expansion, producing a final focused suite of 38 cases.

Coverage area	Cases retained	Reason
Tier limits	TC_001-004, TC_007-010, TC_013-017	Direct validation of tier transaction limits, cumulative daily limits and both supported transfer destinations.
Midnight reset	TC_019, TC_020, TC_022	Directly validates the daily reset rule and the unresolved timezone question.
Pending, failed and reversed	TC_025-028, TC_031	Directly validates how transfer states consume or restore allowance.
Own linked accounts	TC_036-038	Validates that exempt transfers do not consume normal allowance and that the exemption is restricted to the user's own linked account.
Scheduled transfers	TC_043, TC_052	Validates future-date scheduling and single execution at the stated time.
Immediate KYC upgrades	TC_054-056	Validates immediate use of the upgraded tier without re-login.
Remaining-limit message	TC_061, TC_063	Validates rejection and the numerical remaining allowance shown to the user.
Backend enforcement	TC_065	Checks that the stated limits remain authoritative when requests reach the backend.

WHY THESE CASES WERE KEPT

Each retained case either validates an explicit acceptance criterion or tests a necessary boundary of a stated numerical rule. None requires an invented feature such as cancellation, downgrade, partial reversal or audit logging.

5. Cases Retained with Corrections

Case	Correction made
TC_032	A Day 1 pending transfer that fails on Day 2 must not increase Day 2 allowance. Only the original Day 1 usage record is corrected.
TC_033	A Day 1 transfer reversed on Day 2 must not add allowance to Day 2. Its status and original-period accounting are updated.
TC_035	Expanded the own-account exemption check across Tier 1, Tier 2 and Tier 3.
TC_042	Expanded into the end-to-end scheduled-transfer case: creation for a future date, no immediate execution, and execution at 8:00 AM.
TC_047	Removed the invented reservation rule. Limits are evaluated on the scheduled execution date; insufficient allowance causes failure with the remaining-limit error.
TC_058	Changed to clarification-dependent. The specification does not state how earlier same-day usage is treated after an upgrade.

6. Redundant or Merged Cases

These 18 cases were removed or merged because they repeated stronger cases or fragmented one behaviour into unnecessary variations.

Case	Decision and reason
TC_005	Repeats the remaining-limit rejection covered more clearly by TC_061.
TC_006	Repeats TC_004's exact Tier 1 daily-limit boundary.
TC_011	Repeats the remaining-limit scenario for Tier 2.
TC_012	Repeats TC_010's exact Tier 2 daily-limit boundary.
TC_021	Repeats TC_019 and TC_020 midnight behaviour and adds an implementation-level ledger check.
TC_024	Repeats consecutive-day and timezone coverage in TC_019, TC_020 and TC_022.
TC_041	Repeats TC_036: exempt transfers must not consume normal allowance.
TC_044	Overlaps TC_042 and TC_043; scheduling for "today" is not defined.
TC_045	Redundant once the end-to-end scheduled-date case verifies no early execution.
TC_046	Merged into TC_042's 8:00 AM execution check.
TC_049	Merged into corrected TC_047's execution-day allowance evaluation.
TC_053	Repeats TC_052 and adds unsupported performance expectations.
TC_057	Repeats immediate-upgrade checks in TC_054-056.
TC_062	Merged into TC_061; a separate API contract was not supplied.
TC_066	Repeats concurrency and limit-integrity concerns outside the acceptance criteria.
TC_067	Repeats duplicate-execution protection and extends it without specification support.
TC_069	Repeats invalid amount-format variants not defined by the specification.
TC_070	Repeats failure handling and adds insufficient balance, which is outside this feature.

7. Unsupported or Hallucinated Cases

These 20 cases were removed because the AI converted possible system concerns into confirmed requirements or invented capabilities that the specification never provided.

Case	Reason removed
TC_018	Currency precision, rounding and supported currencies were not specified.
TC_023	Daylight-saving behaviour was not stated; it depends on a timezone decision.
TC_029	Duplicate callbacks and callback idempotency assume an event-driven implementation.
TC_030	Duplicate reversal events were not specified.
TC_034	Partial reversals invent a product capability.
TC_039	Unlinking accounts, ownership changes and caching were not defined.
TC_040	API interception and server-side ownership resolution assume internal design.
TC_048	Insufficient wallet balance is a general transfer rule, not part of this feature spec.
TC_050	Scheduled-transfer cancellation was never stated.
TC_051	Scheduled-transfer editing was never stated.
TC_059	Simultaneous KYC completion and transfer submission was unrealistic and unspecified.
TC_060	Tier downgrades and KYC revocation were not part of the requirements.
TC_064	Locale, accessibility and currency-symbol rules were not supplied.
TC_068	Zero and negative amounts belong to general amount validation, not this limits specification.
TC_071	Authentication and authorisation were outside the supplied feature scope.
TC_072	Post-authorisation payload tampering was unrelated to the seven requirements.
TC_073	Database rollback and failure injection assume internal architecture.
TC_074	Processor timeouts and unknown downstream outcomes were not defined.
TC_075	Audit fields, correlation IDs and logging requirements were not stated.
TC_076	Rate limiting, monitoring and abusive traffic were outside the feature scope.

8. Detailed Correction Examples

8.1 Cross-day failure: TC_032

AI-GENERATED PROBLEM

The original expected result said restoration should be applied to the “correct limit period according to an agreed rule.” This avoided defining the outcome and could wrongly increase the current day’s allowance.

Human correction

- The transfer was pending and counted on Day 1.
- The limit resets at midnight.
- If the Day 1 transfer fails on Day 2, Day 2 allowance remains unchanged.
- The amount is not added to the new day’s allowance, and the allowance never exceeds the tier maximum.

Why the correction matters

If a Tier 1 user has 40,000 remaining on Day 2, a failed Day 1 transfer must not increase the Day 2 allowance above 40,000. Restoration applies only to the same daily period in which the transaction originally consumed allowance.

8.2 Own-account exemption: TC_035

The generated case checked only Tier 1. Human review expanded it across all tiers because the exemption says “all limits” and is not tier-specific.

8.3 Scheduled-transfer evaluation: TC_047

The AI invented a choice between reservation at scheduling time and validation at execution time. Human review corrected the case: creating a schedule does not reserve a future day’s limit. At 8:00 AM on the scheduled date, the current tier and that day’s remaining allowance are checked. If allowance is insufficient, execution fails and the error shows the remaining limit.

8.4 KYC upgrade calculation: TC_058

The AI assumed earlier same-day usage carries forward under the new tier. That may be reasonable, but it is not stated. The case was retained only as a clarification-dependent test instead of presenting an assumed numerical result as fact.

9. Where AI Helped

Strength	Observed value
Rapid breadth	AI quickly mapped the seven requirements into limits, reset, status, exemption, scheduling, KYC and rejection scenarios.
Boundary discovery	It proposed exact-limit, below-limit and above-limit values across the three tiers.
Edge-case prompts	It highlighted midnight transitions, pending transfers, reversals and execution-day limit changes that deserved human analysis.
Consistent structure	It produced repeatable IDs, priorities, preconditions, data, steps and expected results, reducing documentation effort.
Review stimulus	Even incorrect cases exposed unanswered questions and helped the tester challenge the specification.

10. Where AI Harmed Test Design

Weakness	Observed risk
Assumptions became requirements	The model treated callbacks, idempotency, database rollback, logging and rate limiting as if they were specified.
Invented capabilities	It added cancellation, editing, partial reversals, downgrades and other behaviours absent from the feature.
Duplicate coverage	It repeated the same limit, reset, exemption and scheduling logic under multiple IDs, making the suite look more complete than it was.
Incorrect business logic	The cross-day restoration case could have increased the wrong day's allowance.
False certainty	Clarification-dependent behaviour was sometimes written as a definite expected result.
Scope dilution	Generic security and transfer-validation scenarios distracted from traceability to the seven acceptance criteria.

KEY LESSON

AI improved speed and breadth, but without human review it produced a larger and less defensible suite. Test count was not a reliable measure of test quality.

11. Final Evaluation

Measure	Result
Initial AI-generated cases	76
Retained without major correction	32
Retained after correction or expansion	6
Redundant or merged cases	18
Unsupported or hallucinated cases	20
Final focused suite	38

12. Conclusion

The exercise demonstrated that AI is effective as a test-design assistant but unreliable as the final authority. It accelerated discovery, organised scenarios and surfaced useful boundaries. However, it also duplicated coverage, invented product capabilities, assumed internal architecture and produced incorrect cross-day business logic.

Human QA judgement was required to trace each case to the specification, distinguish useful questions from executable tests, correct expected results and remove unsupported content. The strongest outcome came from combining AI-generated breadth with manual specification analysis and domain reasoning.

FINAL POSITION

AI should propose and challenge. The QA engineer must validate, prioritise and take responsibility for every test submitted.

Review declaration

I manually reviewed the generated cases against the supplied specification, challenged unsupported assumptions, corrected the business logic and can explain the final classifications and decisions in a live discussion.